

# Kubernetes discovery chaotic cache and how to speed up your Velero backups?

**DORS/CLUC**

**29<sup>TH</sup>** OPEN SYSTEMS DAYS  
CROATIAN LINUX USERS' CONFERENCE


**16 -17 MAY 2024**

Zagreb - Visoko učilište Algebra, Gradišćanska 24

Dario Sindičić

# Who are we?

- Government IT company with over 50 years of existence
- Ministry of Finance, Tax administration, Customs Administration, Elections (local, presidential, parliamentary)
- In the current transition into the mostly containerized environment
- And those new stuff also needs backup



# Introduction



```
kind: Namespace
apiVersion: v1
metadata:
  name: a00-dors-cluc-dev
spec:
  finalizers:
    - kubernetes
```

```
---
kind: ConfigMap
apiVersion: v1
metadata:
  name: oib-config
  namespace: a00-dors-cluc-dev
immutable: false
data:
  DATABASE: database01.something.com
```

```
kind: Secret
apiVersion: v1
metadata:
  name: database-cred
  namespace: a00-dors-cluc-dev
data:
  DB_PASS: b3B1bnNlc2FtZQ==
  DB_USR: YWRtaW4=
type: Opaque
```

```
apiVersion: v1
kind: Pod
metadata:
  name: oib-fetch
  namespace: a00-dors-cluc-dev
spec:
  containers:
    - name: oib-fetcher
      image: 'oib-fetcher:latest'
      ports:
        - containerPort: 8080
      envFrom:
        - configMapRef:
            name: oib-config
          secretRef:
            name: database-cred
```



# Our project

**NS** a00-dors-cluc-dev

**S** database-cred

**CM** oib-config

**D** oib-fetcher


**IS** oib-fetcher

**SA** default

Bunch  
of  
YAMLs



B  
a  
c  
k  
u  
p



a00-dors-cluc-dev-backup-17-05-24

**NS** a00-dors-cluc-dev

**S** database-cred

**CM** oib-config

**D** oib-fetcher

**IS** oib-fetcher

**SA** default



# Backup creation

```
sh-4.4$ /velero backup create dors-one \
                                     --include-namespaces a00-dors-cluc-dev
Backup request "dors-one" submitted successfully.
Run `velero backup describe dors-one` or `velero backup logs dors-
one` for more details
```

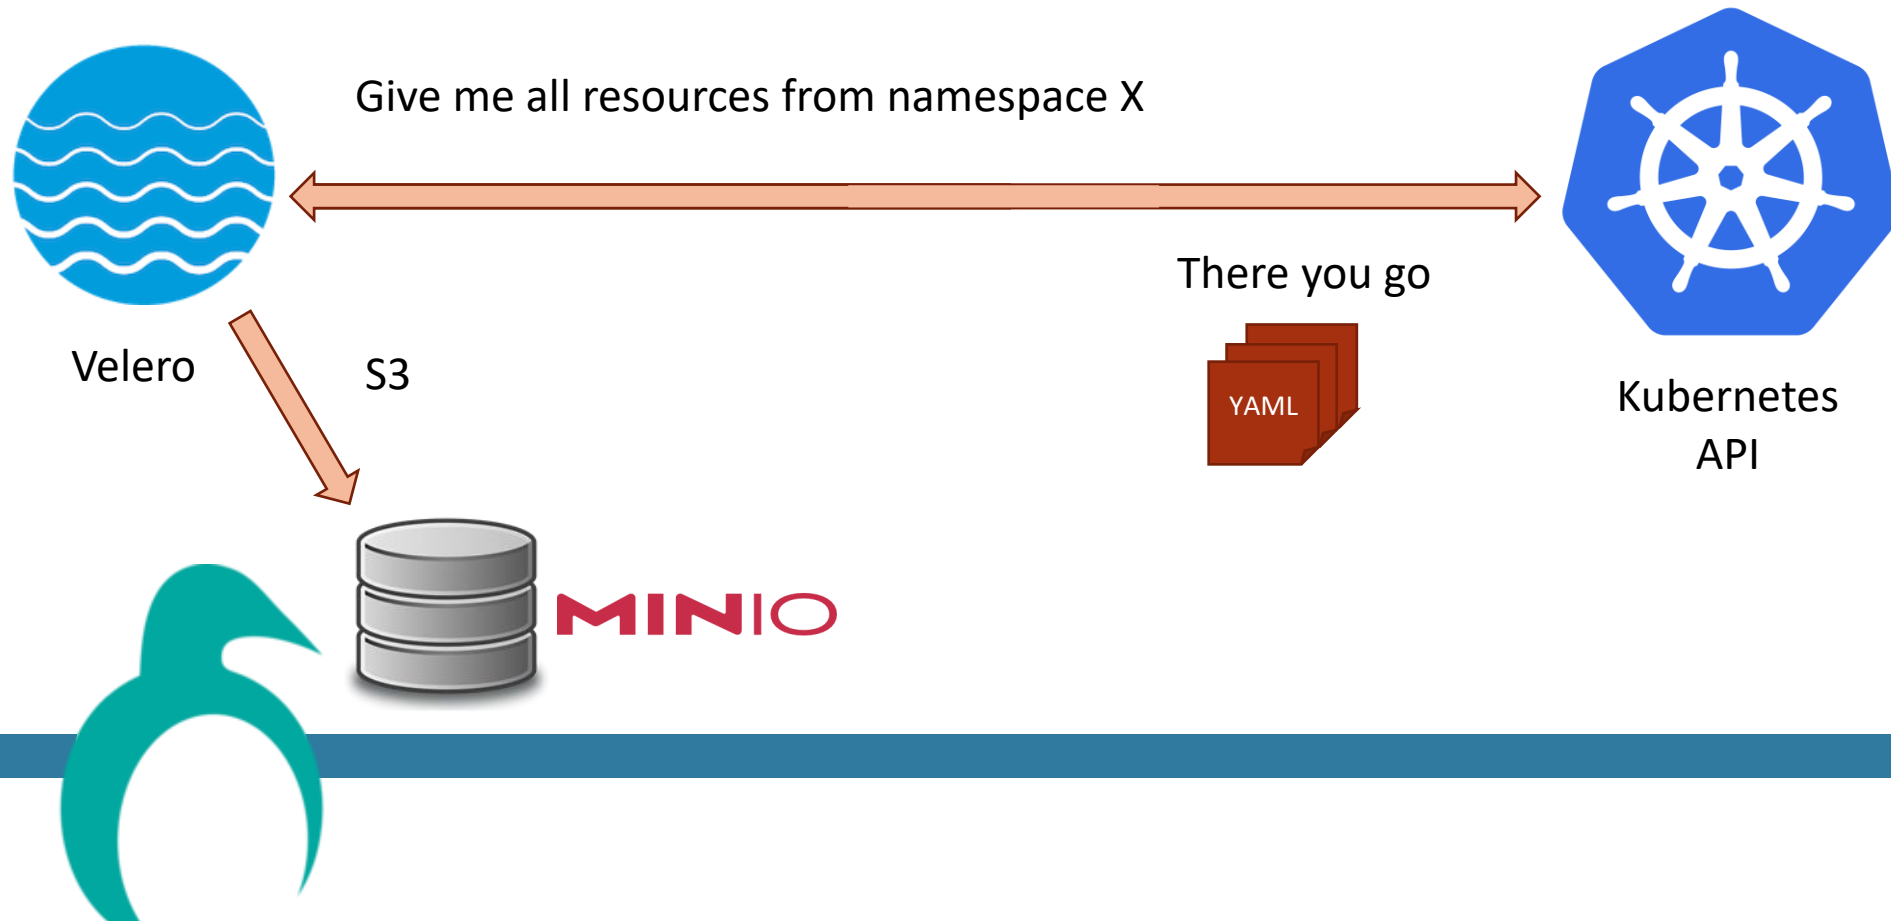


```
apiVersion: velero.io/v1
kind: Backup
metadata:
  name: backup-resources-dors-first-take
spec:
  ...
  includedNamespaces:
    - a00-dors-cluc-dev
  ...
status:
  completionTimestamp: '2024-05-04T15:40:21Z'
  phase: Completed
  progress:
    itemsBackedUp: 45
    totalItems: 45
  startTimestamp: '2024-05-04T15:38:46Z'
```

- small project without any volume
- Total time = 95 seconds
- 412 projects in production → 10 hours and 38 minutes minimal



# Architecture



# Disk

## Velero log: Backup is in memory

```
time="2024-05-05T12:00:45Z" level=info msg="Backed up 45 items out of an estimated total of 45
```

## Strace from Velero: TLS handshake ➔ saving data on Minio

```
12:00:45 connect(12, {sa_family=AF_INET, sin_port=htons(9000), sin_addr=inet_addr("172.30.62.27")})  
12:00:45 write(12,"_26\3\1\1.\.....n\300\24\0\234\0\235\0\0005\300\22\0\n\23\1\23\2\23\3\1\0\0\273\0\0\0.\0\0\0)guardian-minio-svc.apis-backup-restore.svc\0\5\0\5\1\0\0\0\0\0\n\0\n\0\10\0\35\0\27\0\30\0\31\0\0\0\2\1\0\0\0\r\0\32\0\30\0\10\.....", 307) = 307 <0.002225>
```

## Velero log: Backup is saved on Minio

```
time="2024-05-05T12:00:46Z" level=info msg="Backed up 45 items out of an estimated total of 45
```





# Architecture




```
9706 request.go:601] waited for 21.595097708s due to client-side  
throttling, not priority and fairness, request:  
GET:https://172.30.48.1:443/apis/application.isf.com
```

```
9706 request.go:601] waited for 11.395595197s due to client-side  
throttling, not priority and fairness, request:  
GET:https://172.30.48.1:443/apis/config.io/v1
```



# Client-side throttling



```
const (  
    DefaultQPS    float32 = 5.0  
    DefaultBurst int      = 10  
)
```



client-go/rest/config.go



```
sh-4.4$ /velero server --help
Run the velero server
```

Usage:

```
velero server [flags]
```

Flags:

<code>--backup-sync-period duration</code>	How often to ensure all Velero backups in object storage exist as Backup API
if none is explicitly specified for a backup storage location. (default 1m0s)	
<code>--client-burst int</code>	Maximum number of requests by the server to the Kubernetes API in a short per
<code>--client-page-size int</code>	Page size of requests by the server to the Kubernetes API when listing objec
500)	
<code>--client-qps float32</code>	Maximum number of requests per second by the server to the Kubernetes API on

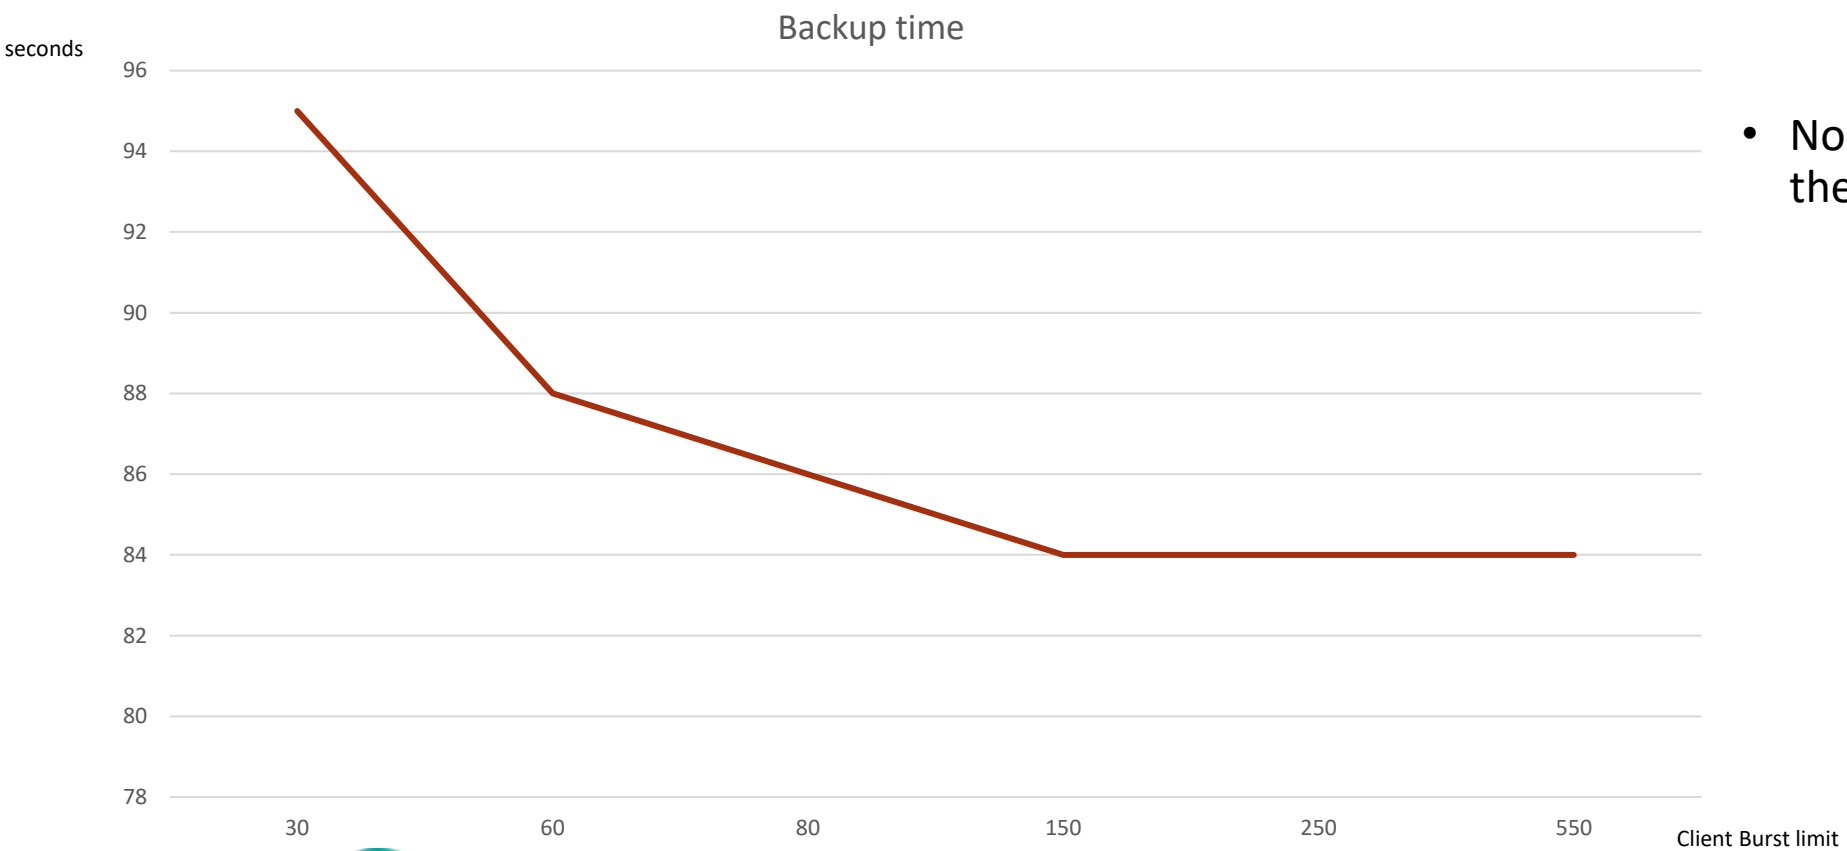
pkg/cmd/server/server.go



```
f.SetClientQPS(config.clientQPS)

if config.clientBurst <= 0 {
    return nil, errors.New("client-burst must be positive")
}
f.SetClientBurst(config.clientBurst)
```





- No significant changes in the backup speed



```
9706 request.go:601] waited for 21.595097708s due to client-side  
throttling, not priority and fairness, request:  
GET:https://172.30.48.1:443/apis/application.isf.com
```

```
9706 request.go:601] waited for 11.395595197s due to client-side  
throttling, not priority and fairness, request:  
GET:https://172.30.48.1:443/apis/config.io/v1
```



# Architecture





# CUSTOM RESOURCE

```
kind: Namespace
apiVersion: v1
metadata:
  name: a00-dors-cluc-dev
spec:
  finalizers:
    - kubernetes
```

```
---
kind: ConfigMap
apiVersion: v1
metadata:
  name: oib-config
  namespace: a00-dors-cluc-dev
immutable: false
data:
  DATABASE: database01.something.com
```

```
kind: DorsCluc
apiVersion: v1
metadata:
  name: dors
  namespace: cluc
spec:
  events:
    - year: 2024
      location: Algreba
    - year: 2023
      location: FER
    - year: 2013
      location: Hotel International
```

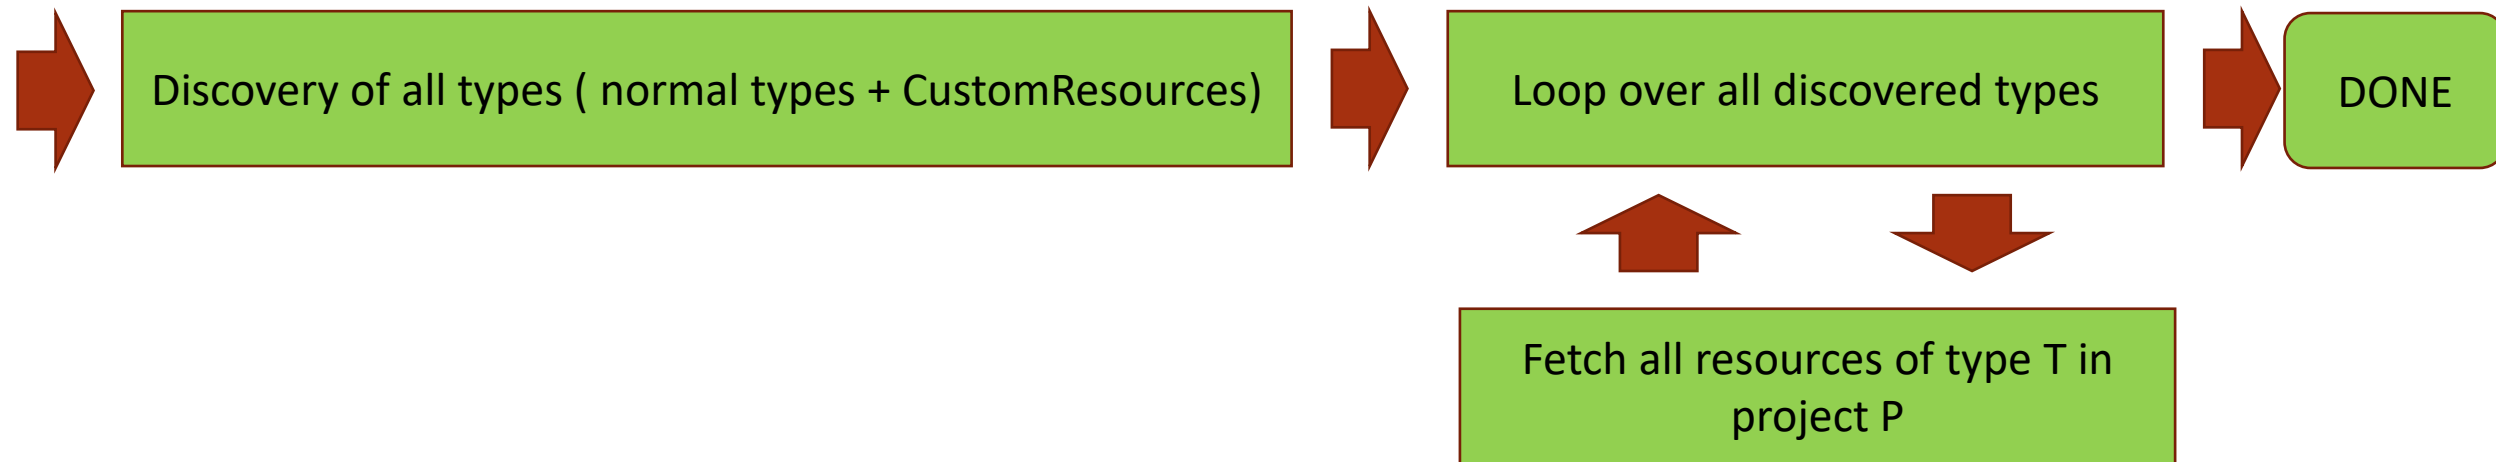
```
v1
fetch
  a00-dors-cluc-dev
:
oib-fetcher
  'oib-fetcher:latest'
containerPort: 8080
envFrom:
  - configMapRef:
      name: oib-config
    secretRef:
      name: database-cred
```



# Velero steps



VELERO




# Analysis of API calls


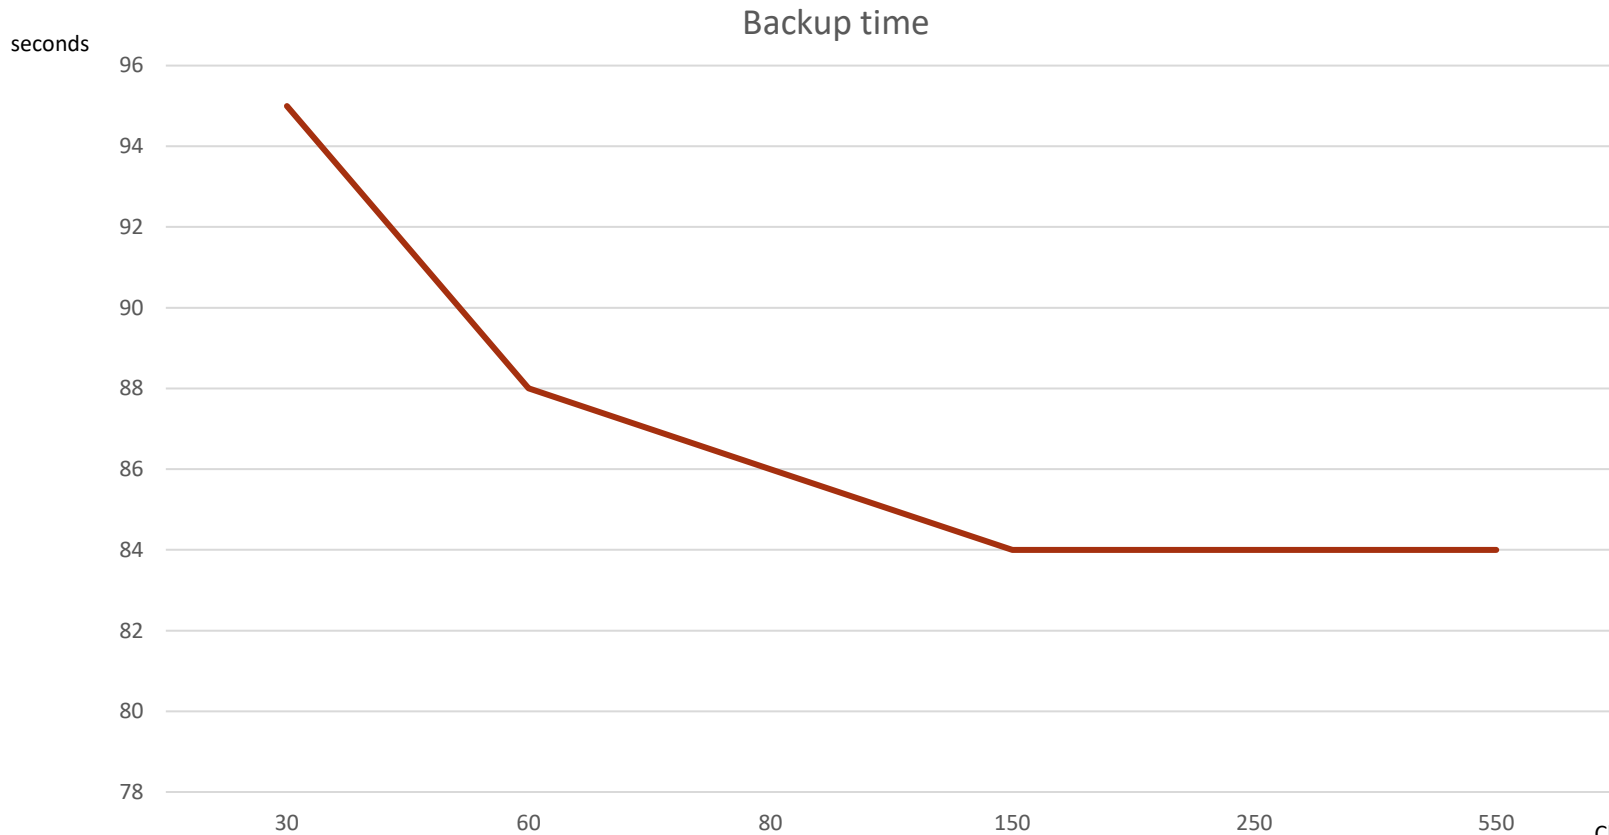
Total number of API calls: 284

Number of API calls related to the backup      82      187      15



82 seconds of doing "something"





- 284 calls
- 30 r/s → 284 calls → 10 sec
- 60 r/s → 284 calls → 4.5 sec
- 150 r/s → 284 calls → 2 sec
- 250 r/s → 284 calls → 1.5 sec



# Doing „something“

waited for 21.595097708s due to client-side throttling

- Probably client-side throttling are caused by discovery of all types
- Let's fork the code and see what to do with it

<https://github.com/vmware-tanzu/velero>



```
func (h *helper) Refresh() error {

    b := make([]byte, 2048) // adjust buffer size to be larger than expected stack
    n := runtime.Stack(b, false)
    s := string(b[:n])

    t1 := time.Now().UnixNano()

    h.logger.Info("Dario Running the velero refresh DiscoveryCache.")
    h.logger.Infof("Dario Running the velero refresh DiscoveryCache StackTrace.", s)

    h.lock.Lock()
    defer h.lock.Unlock()

    groupResources, err := restmapper.GetAPIGroupResources(h.discoveryClient)
    if err != nil {
        return errors.WithStack(err)
    }

    var serverResources []*metav1.APIResourceList
```

```
}

h.serverVersion = serverVersion

    t2 := time.Now().UnixNano()
    h.logger.Infof("Dario Completed the velero refresh DiscoveryCache StackTrace.", s)
    h.logger.Infof("Dario Completed the velero time.", (t2 - t1) / 1000000000.0)

    return nil
}
```

pkg/discovery/helper.go



- Let's add some debug prints to see what is going on

14:20:54Z" level=info msg="Dario Running the velero refresh  
DiscoveryCache.

14:21:35Z" level=info msg="Dario Completed the velero time 41 seconds"

14:21:35Z" level=info msg="Dario Running the velero refresh  
DiscoveryCache.

14:22:16Z" level=info msg="Dario Completed the velero time 41 seconds"

- 2 discovery searches  
was run by the Velero

**82 seconds of doing "something",  
41 seconds each**

- 457 api resources =>  
~ 10 qps



# Plugin

```
// TODO(ncdc): consider a k8s style WantsKubernetesClientSet initialization approach
clientset, err := f.KubeClient()
if err != nil {
    return nil, err
}

discoveryHelper, err := velerodiscovery.NewHelper(clientset.Discovery(), logger)
if err != nil {
    return nil, err
}
```

```
clientset, err := f.KubeClient()
if err != nil {
    return nil, err
}

logger.Info("Creating new helper in newRemap CRD Version action plugin")
discoveryHelper, err := velerodiscovery.NewHelper(clientset.Discovery(), logger)
```

pkg/cmd/server/plugin/plugin.go



newServiceAccountBackupItemAction function

newRemapCRDVersionAction function

- Plugin is ran in seperate process so this factory has different settings than default Velero one
- It is using really low QPS and burst parameters





```
// NewFactory returns a Factory.
func NewFactory(baseName string, config VeleroConfig) Factory {
    f := &factory{
        flags:    pflag.NewFlagSet("", pflag.ContinueOnError),
        baseName: baseName,
    }
    f.clientQPS = 300.0;
    f.clientBurst = 600;
    f.namespace = os.Getenv("VELERO_NAMESPACE")
    if config.Namespace() != "" {
        f.namespace = config.Namespace()
    }

    // We didn't get the namespace via env var or config file, so use the default.
    // Command line flags will override when BindFlags is called.
    if f.namespace == "" {
        f.namespace = veleroapi.DefaultNamespace
    }
}
```





# End results

- Backup time reduced to only 2 seconds
- Total backup of production with only one instance of Velero = 14 minutes

```
apiVersion: velero.io/v1
kind: Backup
metadata:
  name: backup-resources-dors-n-th-take
spec:
  ...
  includedNamespaces:
    - a00-dors-cluc-dev
  ...
status:
  completionTimestamp: '2024-05-10T16:49:46Z'
  phase: Completed
  progress:
    itemsBackedUp: 45
    totalItems: 45
  startTimestamp: '2024-05-10T16:49:44Z'
```





Thank you !



# Appendix

```
// initDiscoveryHelper instantiates the server's discovery helper and spawns a
// goroutine to call Refresh() every 5 minutes.
func (s *server) initDiscoveryHelper() error {
    discoveryHelper, err := velerodiscovery.NewHelper(s.discoveryClient, s.logger)
    if err != nil {
        return err
    }
    s.discoveryHelper = discoveryHelper

    go wait.Until(
        func() {
            if err := discoveryHelper.Refresh(); err != nil {
                s.logger.WithError(err).Error("Error refreshing discovery")
            }
        },
        5*time.Minute,
        s.ctx.Done(),
    )

    return nil
}
```

```
// getAllItems gets all relevant items from all API groups.
func (r *itemCollector) getAllItems() []*kubernetesResource {
    var resources []*kubernetesResource
    for _, group := range r.discoveryHelper.Resources() {
        groupItems, err := r.getGroupItems(r.log, group)
        if err != nil {
            r.log.WithError(err).WithField("apiGroup", group.String()).Error(err)
            continue
        }

        resources = append(resources, groupItems...)
    }

    return resources
}
```



```
if config.clientBurst <= 0 {
    return nil, errors.New("client-burst must be positive")
}
f.SetClientBurst(config.clientBurst)

if config.clientPageSize < 0 {
    return nil, errors.New("client-page-size must not be negative")
}

kubeClient, err := f.KubeClient()
if err != nil {
    return nil, err
}

veleroClient, err := f.Client()
```

```
backupper, err := backup.NewKubernetesBackupper(
    s.veleroClient.VeleroV1(),
    s.discoveryHelper,
    client.NewDynamicFactory(s.dynamicClient),
    podexec.NewPodCommandExecutor(s.kubeClientConfig, s.kubeClient.CoreV1().RESTClient()),
    s.resticManager,
    s.config.podVolumeOperationTimeout,
    s.config.defaultVolumesToRestic,
    s.config.clientPageSize,
)
```

